





5-channel policy input tensor (after 5 ADMM iterations — first state the policy evaluates)



Ch 0: x (recon)
range [0.06, 0.82]




Ch 1: z (data-consistent)
range [0.07, 0.82]



Ch 2: u (accumulated residual)
range [-0.050, 0.032]



Ch 3: noise_level = 0.050
(measurement noise std)



Ch 4: iter_frac = 0.25
(5 of ~20 iters)

